To the Editors,
International Journal of Heat and Mass Transfer

Dear Editors,

Please consider our manuscript, "The Boussinesq–compressible gap in steady two-dimensional buoyancy-driven enclosure convection is linear to leading order in the relative temperature contrast: a factorised physics prior for surrogates", for publication in the *International Journal of Heat and Mass Transfer* as a full-length article.

Surrogates for thermal convection are trained overwhelmingly on Boussinesq data, yet the flows engineers need are compressible. The two are close in integral quantities such as the Nusselt number and differ in the local fields that surrogates are asked to reproduce. The manuscript establishes what sets the size of that difference and turns the answer into a design rule.

- **One axis.** In the non-dimensionalisation used by the released paired data and by our solver, the Froude number is pinned ($Fr^2 = \varepsilon$ identically) and the Mach number enters at $O(Ma^2)$; a 24-fold change in Ma leaves the measured correction unchanged at the precision of the solve. The break-symmetry group is therefore one-dimensional, spanned by $\varepsilon = \Delta T/T_0$.
- **A measured law.** Across 627 paired two-dimensional solves in the unit square at $Pr = 0.71$ spanning $\varepsilon = 0.02$–1.9, the relative correction obeys $R = A\varepsilon^m$ with $m = 1.02$ in u_x, u_y and p (and the same exponent, with larger scatter, in θ), a vanishing intercept, and an ε^2 curvature that matches the asymptotic expansion. In 104-sample checks in a 2:1 geometry and at $Pr = 4.38$ the exponent is statistically unchanged; the amplitude reproduces on an independently generated public data set at $\varepsilon = 1.2$.
- **A design rule.** Because the correction factorises as $D = \varepsilon P(X)$, a network that predicts only P has a 1.09× lower channel-averaged error than the strongest conditional baseline in extrapolation (θ and p individually; the u_x/u_y intervals include zero), and 1.89× better than an ε-blind network when ε is withheld, which makes single-contrast public data reusable at other contrasts within one boundary-condition family, without retraining.
- **Wall heat transfer.** At $\varepsilon = 1.2$ the integral Nusselt number changes by 0.6%, while the local wall heat flux is redistributed by 4.25%; the integral agreement therefore conceals a local error that a surrogate must be trained to resolve.

IJHMT has published numerical methodology for non-Oberbeck–Boussinesq flows (Demou, Frantzis and Grigoriadis, *Int. J. Heat Mass Transfer* 125 (2018) 1156–1168), so the present study fits the journal's scope: it quantifies that effect and turns it into a surrogate design rule.

The main sweep is two-dimensional and steady; a preliminary twelve-sample three-dimensional check reproduces the exponent, and a POD-ROM with at most four modes on the same samples lowers the mean u_x error from 0.118 to 0.054 over four held-out contrasts. We state the remaining limits (one Rayleigh number, $res = 16$, $\varepsilon \leq 1.2$) explicitly. Every number in the paper is regenerated from stored solver output by the scripts in the repository; the public data set of the companion study is used unchanged, and our scripts and machine-readable results are released with the manuscript; the raw solver output is available from the authors on request.

The work is not under consideration elsewhere, and the authors declare no competing interests.

Yours sincerely,
[Corresponding author name]
[Affiliation]
[Email]